

## Workflow

1. House Price dataset
2. Data preprocessing
3. Data Analysis
4. Train and Test Splitting
5. XGBoost Regression Algorithm
6. Evaluation (By Testing Dataset)

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import sklearn.datasets
from sklearn.model_selection import train_test_split
from xgboost import XGBRegressor
from sklearn import metrics

# importing Boston House Price Dataset
house_price_dataset = sklearn.datasets.fetch_california_housing()

# printing the california house price dataset
print(house_price_dataset)
```



```
name': None, 'target_names': ['MedHouseVal'], 'feature_names': ['MedInc', 'HouseAge', 'AveRooms', 'AveBedrms', 'Population', 'AveOccup', 'Latitude', 'Longitude']
```



```
# Loading the dataset to pandas dataframe
house_price_dataframe = pd.DataFrame(house_price_dataset.data)
house_price_dataframe.head()
```



	0	1	2	3	4	5	6	7	
0	8.3252	41.0	6.984127	1.023810	322.0	2.555556	37.88	-122.23	
1	8.3014	21.0	6.238137	0.971880	2401.0	2.109842	37.86	-122.22	
2	7.2574	52.0	8.288136	1.073446	496.0	2.802260	37.85	-122.24	
3	5.6431	52.0	5.817352	1.073059	558.0	2.547945	37.85	-122.25	
4	3.8462	52.0	6.281853	1.081081	565.0	2.181467	37.85	-122.25	



Next steps:

[Generate code with house\\_price\\_dataframe](#)

[View recommended plots](#)

[New interactive sheet](#)

Suggested code may be subject to a license | [showMiko/Projects-Hacktober](#)

```
# if we use "house_price_dataframe = pd.DataFrame(house_price_dataset.data)" we cant get the columns names to get the name use
house_price_dataframe = pd.DataFrame(house_price_dataset.data, columns = house_price_dataset.feature_names)
print(house_price_dataframe)
```



	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup	Latitude	Longitude
0	8.3252	41.0	6.984127	1.023810	322.0	2.555556	37.88	-122.23
1	8.3014	21.0	6.238137	0.971880	2401.0	2.109842	37.86	-122.22
2	7.2574	52.0	8.288136	1.073446	496.0	2.802260	37.85	-122.24
3	5.6431	52.0	5.817352	1.073059	558.0	2.547945	37.85	-122.25
4	3.8462	52.0	6.281853	1.081081	565.0	2.181467	37.85	-122.25
...	...	...	...	...	...	...	...	...
20635	1.5603	25.0	5.045455	1.133333	845.0	2.560606	39.48	-122.43
20636	2.5568	18.0	6.114035	1.315789	356.0	3.122807	39.49	-122.43
20637	1.7000	17.0	5.205543	1.120092	1007.0	2.325635	39.43	-122.43
20638	1.8672	18.0	5.329513	1.171920	741.0	2.123209	39.43	-122.43
20639	2.3886	16.0	5.254717	1.162264	1387.0	2.616981	39.37	-122.43

Longitude

```

0      -122.23
1      -122.22
2      -122.24
3      -122.25
4      -122.25
...      ...
20635   -121.09
20636   -121.21
20637   -121.22
20638   -121.32
20639   -121.24


```

[20640 rows x 8 columns]



```

# here price of the column as not included bcz we just included the features name to include price just import the target column
house_price_dataframe['price'] = house_price_dataset.target
house_price_dataframe.head()

```



	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup	Latitude	Longitude	price
0	8.3252	41.0	6.984127	1.023810	322.0	2.555556	37.88	-122.23	4.526
1	8.3014	21.0	6.238137	0.971880	2401.0	2.109842	37.86	-122.22	3.585
2	7.2574	52.0	8.288136	1.073446	496.0	2.802260	37.85	-122.24	3.521
3	5.6431	52.0	5.817352	1.073059	558.0	2.547945	37.85	-122.25	3.413
4	3.8462	52.0	6.281853	1.081081	565.0	2.181467	37.85	-122.25	3.422

Next steps: [Generate code with house\\_price\\_dataframe](#) [View recommended plots](#) [New interactive sheet](#)

```

# checking the numbers of rows and columns in our dataset
house_price_dataframe.shape

```

```

(20640, 9)

```

```

# handling missing values
house_price_dataframe.isnull().sum()

```



	0
MedInc	0
HouseAge	0
AveRooms	0
AveBedrms	0
Population	0
AveOccup	0
Latitude	0
Longitude	0
price	0

## Stastical Measures of the Dataset

1. Count - No of data points we have (rows).
2. Mean - value for the respective column(Average).
3. std - standard deviation.

```

# stastical measures of the dataset
house_price_dataframe.describe()

```

	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup	Latitude	Longitude	price
count	20640.000000	20640.000000	20640.000000	20640.000000	20640.000000	20640.000000	20640.000000	20640.000000	20640.000000
mean	3.870671	28.639486	5.429000	1.096675	1425.476744	3.070655	35.631861	-119.569704	2.068558
std	1.899822	12.585558	2.474173	0.473911	1132.462122	10.386050	2.135952	2.003532	1.153956
min	0.499900	1.000000	0.846154	0.333333	3.000000	0.692308	32.540000	-124.350000	0.149990
25%	2.563400	18.000000	4.440716	1.006079	787.000000	2.429741	33.930000	-121.800000	1.196000
50%	3.534800	29.000000	5.229129	1.048780	1166.000000	2.818116	34.260000	-118.490000	1.797000
75%	4.743250	37.000000	6.052381	1.099526	1725.000000	3.282261	37.710000	-118.010000	2.647250
max	15.000100	52.000000	141.909091	34.066667	35682.000000	1243.333333	41.950000	-114.310000	5.000010

## ✓ Check for the correlation of the dataset

1. +ve - correlation. (one inc other also inc.)
2. -ve - correlation. (one dec other also dec.)


Suggested code may be subject to a license | [ajaydurgam85/HousePrice\\_Prediction\\_2](#)  
`correlation = house_price_dataframe.corr()`

## ✓ commands explanation

1. `fmt='.1f'` = 0.0, 0.1 (if `fmt='.2f'` = 0.00, 0.11).
2. `square=True` plot in square shape.
3. `cbar =True` for a bar in the right side.
4. `annot =` to print the numbers inside the box.
5. `annot_kws =` for sizes of the numbers inside the box.
6. `cmap` for color of the map

```
plt.figure(figsize = (8,8))
sns.heatmap(correlation, cbar =True, square=True, fmt='.2f', annot = True, annot_kws ={'size':8}, cmap='Reds' )
```

<Axes: >



```
# splitting the features and Target
x = house_price_dataframe.drop(['price'], axis=1) # axis 1 for representing the columns
y = house_price_dataframe['price']
```

```
print(x)
print(y)
```

```
↩
MedInc  HouseAge  AveRooms  AveBedrms  Population  AveOccup  Latitude \
0      8.3252    41.0    6.984127  1.023810      322.0    2.555556    37.88
1      8.3014    21.0    6.238137  0.971880     2401.0    2.109842    37.86
2      7.2574    52.0    8.288136  1.073446      496.0    2.802260    37.85
3      5.6431    52.0    5.817352  1.073059      558.0    2.547945    37.85
4      3.8462    52.0    6.281853  1.081081      565.0    2.181467    37.85
...      ...      ...      ...      ...      ...      ...      ...
20635  1.5603    25.0    5.045455  1.133333      845.0    2.560606    39.48
20636  2.5568    18.0    6.114035  1.315789      356.0    3.122807    39.49
20637  1.7000    17.0    5.205543  1.120092     1007.0    2.325635    39.43
20638  1.8672    18.0    5.329513  1.171920      741.0    2.123209    39.43
20639  2.3886    16.0    5.254717  1.162264     1387.0    2.616981    39.37
```

```
Longitude
0      -122.23
1      -122.22
2      -122.24
3      -122.25
4      -122.25
...      ...
20635  -121.09
20636  -121.21
20637  -121.22
20638  -121.32
20639  -121.24
```

```
[20640 rows x 8 columns]
```

```
0      4.526
1      3.585
2      3.521
3      3.413
4      3.422
...
20635  0.781
20636  0.771
20637  0.923
20638  0.847
20639  0.894
```

```
Name: price, Length: 20640, dtype: float64
```

```
# splitting the dataset for testing and training
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=2)
print(x.shape, x_train.shape, x_test.shape)
```

```
↩ (20640, 8) (16512, 8) (4128, 8)
```

## ✓ Model Training

### XGBoost Regressor

1. It is an Ensemble model
2. this basically incorporate the one or two model to achieve the result.

```
model = XGBRegressor()
```

```
model.fit(x_train, y_train)
```

```
↩
XGBRegressor
XGBRegressor(base_score=None, booster=None, callbacks=None,
             colsample_bylevel=None, colsample_bynode=None,
             colsample_bytree=None, device=None, early_stopping_rounds=None,
             enable_categorical=False, eval_metric=None, feature_types=None,
             gamma=None, grow_policy=None, importance_type=None,
             interaction_constraints=None, learning_rate=None, max_bin=None,
             max_cat_threshold=None, max_cat_to_onehot=None,
             max_delta_step=None, max_depth=None, max_leaves=None,
             min_child_weight=None, missing=nan, monotone_constraints=None,
             multi_strategy=None, n_estimators=None, n_jobs=None,
             num_parallel_tree=None, random_state=None, ...)
```

```
train_data_prediction = model.predict(x_train)
```

```
print(train_data_prediction)
```

```
→ [0.5523039 3.0850039 0.5835302 ... 1.9204227 1.952873 0.6768683]
```

```
# R squared error
```

```
error_score1 = metrics.r2_score(y_train, train_data_prediction)  
print("R squared error : ", error_score1)
```

```
# Mean square error
```

```
error_score2 = metrics.mean_squared_error(y_train, train_data_prediction)  
print("Mean square error : ", error_score2)
```

```
→ R squared error : 0.943650140819218  
Mean square error : 0.0748112971690747
```

```
# let visually see the results in graphy using plt(Matplotlib)
```

```
plt.scatter(y_train, train_data_prediction)  
plt.xlabel("Actual Price")  
plt.ylabel("Predicted Price")  
plt.title("Actual Price vs Predicted Price")  
plt.show()
```



## ✓ Evaluation discuss

**Mean square error : 0.0748112971690747**

*It does not higher (like 5 or 6) as it is less than 1, means that our model is performing good on training dataset.*

*The result we got is only for Training Data.*

```
test_data_prediction = model.predict(x_test)
```

```
print(test_data_prediction)
```

```
→ [2.8649795 1.790346 0.92074925 ... 1.5385513 0.92647874 2.043316 ]
```

```
# R squared error
```

```
error_score1 = metrics.r2_score(y_test, test_data_prediction)  
print("R squared error : ", error_score1)
```

```
# Mean square error
```

```
error_score2 = metrics.mean_squared_error(y_test, test_data_prediction)  
print("Mean square error : ", error_score2)
```

```
→ R squared error : 0.8338000331788725  
Mean square error : 0.22387540906811954
```

## ✓ Evaluation discuss

**Mean square error : 0.22387540906811954**

*It does not higher (like 5 or 6) as it is less than 1, means that our model is performing good on training dataset.*

The result we got is only for **TEST Data**.

Suggested code may be subject to a license | PSaiRahul/House-Price-Prediction  
# let visually see the results in graphy using plt(Matplotlib)  
plt.scatter(y\_test, test\_data\_prediction)  
plt.xlabel("Actual Price")  
plt.ylabel("Predicted Price")  
plt.title("Actual Price vs Predicted Price")  
plt.show()

